

# Introduction to Go - Final Project

## Distributed Key-Value Store with Raft Consensus

---

### Introduction

Welcome to your final project! This assignment is designed to push your understanding of Go to its limits by building a **Distributed Key-Value Store**. Distributed systems are where Go truly shines due to its first-class support for concurrency, networking, and efficient binary serialization.

In this project, you will move beyond a single-server application and build a cluster of nodes that must agree on data state even if some nodes fail. You will implement a simplified version of a consensus algorithm (like Raft) to ensure data consistency across multiple instances.

This project will challenge you to:

- Implement complex **Network Communication** between peer nodes.
  - Master **Channels and Select** statements for coordination.
  - Use **Custom Serialization** (like GOB or JSON) for log replication.
  - Handle **Partial Failures** and network partitions.
  - Implement a **Distributed Background Job** (log compaction).
- 

### Project Overview

You will create a distributed system where multiple Go processes work together to store data. Each node will provide a REST API for users to interact with, but behind the scenes, the nodes will communicate with each other to replicate data.

The system will consist of:

1. **The Storage Engine:** A thread-safe in-memory map.
  2. **The Consensus Module:** Logic that handles leader election and log replication.
  3. **The Peer Manager:** Maintains a list of active cluster members.
  4. **Log Compaction:** A periodic background task that cleans up old command history to save memory.
- 

### Project Objectives

By completing this project, you will:

- Use **TCP/UDP/RPC** for inter-node communication.
- Implement **State Machines** to manage node roles (Follower, Candidate, Leader).

- Apply **Advanced Concurrency** patterns to manage heartbeat timers and request-vote cycles.
  - Use context.WithTimeout to manage network latency and node unresponsiveness.
  - Implement **Persistent Logs** to allow nodes to recover their state after a crash.
- 

## Project Requirements

### 1. Design Data Models

Define the structures for the cluster state and the replication logs.

#### Node Struct

```
type Node struct {
    ID      int      `json:"id"`
    Address string   `json:"address"` // e.g., "localhost:8080"
    Peers   []string `json:"peers"`
    Role    string   `json:"role"`  // Leader, Follower, Candidate
    Log     []LogEntry `json:"log"`
    CommitIdx int      `json:"commit_index"`
    mu      sync.RWMutex
}
```

#### LogEntry Struct

```
type LogEntry struct {
    Term   int      `json:"term"`
    Command string  `json:"command"` // e.g., "SET key value"
    Key    string   `json:"key"`
    Value  interface{}`json:"value"`
}
```

### 2. Define Interfaces for Replication

Create interfaces to abstract the consensus mechanism.

#### Consensus Interface

```
type Consensus interface {
    RequestVote(term int, candidateID int) (granted bool)
```

```
AppendEntries(term int, leaderID int, entries []LogEntry) error  
StartElection()  
}
```

### 3. Implement RPC Communication

Instead of standard HTTP for everything, use Go's net/rpc package or raw net connections for high-speed node-to-node communication.

- Nodes must "heartbeat" their peers to maintain leadership.
- If a heartbeat is missed, a new election must be triggered concurrently.

### 4. The Client API

Expose a standard RESTful API for the end-user:

- **GET /get/{key}**: Retrieve a value (can be served by any node or forwarded to the leader).
- **POST /set**: Submit a new key-value pair (must be handled by the leader and replicated).
- **GET /cluster/status**: View the current leader and health of all nodes.

### 5. Periodic Log Compaction

To prevent the log from growing infinitely, implement a background goroutine:

- Runs every \$X\$ minutes.
- Identifies entries that have been committed and applied.
- Snapshots the current state and clears the log.

---

## Grading Criteria

Criteria	Points
<b>Replication Logic:</b> Data is successfully copied to at least 2/3 of the nodes.	35
<b>Leader Election:</b> System correctly elects a new leader if the current one dies.	25
<b>Concurrency &amp; Safety:</b> No race conditions during concurrent "SET" operations.	15
<b>Networking:</b> Robust handling of connection timeouts and peer retries.	10

<b>Persistence:</b> Nodes can reload their state from a local JSON/Binary file.	10
<b>Documentation:</b> Clear explanation of how to spin up a 3-node cluster.	5
<b>Total</b>	<b>100</b>

---

## Example Scenario

1. **Node A (Leader)** receives SET name "Gemini".
  2. **Node A** adds the entry to its local log (uncommitted).
  3. **Node A** concurrently sends AppendEntries RPCs to **Node B** and **Node C**.
  4. Once **Node B** responds with success, **Node A** "commits" the entry and responds to the user.
  5. **Node C** (which was offline) comes back online and synchronizes its log automatically with the leader.
- 

## Optional Enhancements

- **Partition Tolerance:** Handle the "Split Brain" scenario where two nodes think they are the leader.
  - **Custom Binary Protocol:** Use encoding/gob instead of JSON for faster replication.
  - **Dynamic Membership:** Implement an endpoint to add a 4th or 5th node to a running cluster.
- 

**Good luck, and happy coding!**